

Introduction

In this project, we derived metadata of 1900 datasets from NYC Open Data using Python, Spark and Pandas. We also extracted more detailed information from selected columns and used different methods to predict their semantic types. In addition, we analyzed the results with the help of applications such as Jupyter Notebook.

General Profiling

Architecture

- The datasets are grouped by sizes (using shell scripts)
 - Small (<50MB), 1825
 - Use small datasets to test logic and correctness
 - Medium(50-100MB), 17
 - Use medium datasets to evaluation performance improvement
 - Large(100MB-1GB), 49
 - Divide spark-submit work evenly to each team member
 - Extra large(>1GB), 9
 - To prevent unforeseen errors and to save time from minor modifications, we store temporary output information after each column is done processing

First Approach

Use RDD

- Extract the header information
- Use map, filter, distinct, count, reduceByKey, sortBy, etc. to extract the number of non-empty cells, empty-cells, distinct values, and top-5 frequent values
- Use type convert to identify if a value is an integer or real(float)
- Use python's dateutil.parser library to identify datetime values
- If a value doesn't belong to integer, real, or datetime, identify it as a text(string) value.

Improvements

- Use Pandas dataframe
- Read the dataset column by column to reduce the memory cost (for large data sets)
- Avoid Python loops, use more Pandas's and NumPy's build-in functions such as max, min, mean, std, head, tail for optimization
 - Also reduce code complexity
- Avoid the use of Python libraries (or even Pandas or NumPy libraries)
 - Try to write our own map functions when possible
- Use Linux Shell scripts and screen command to help create multiple threads
- Evenly split the large datasets to each person, run separately and combine the results afterwards

Results

Task1 Data Type Count

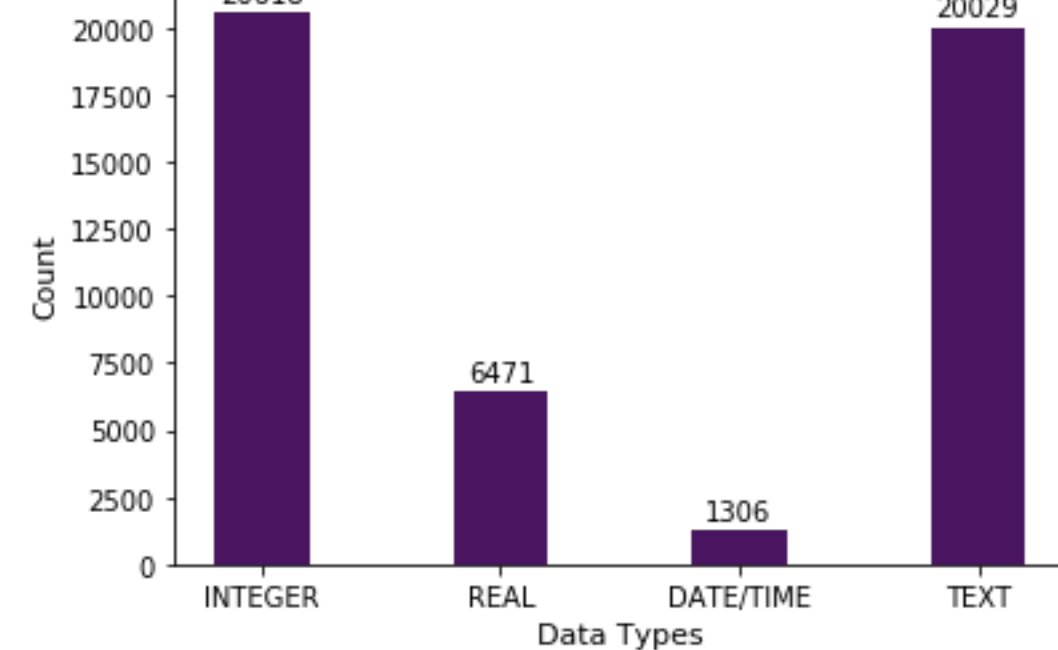


Table. Most Common Types
- Frequent Item Sets

Support	Itemsets
0.521658	(INTEGER)
0.506755	(TEXT)
0.163723	(REAL)
0.140598	(TEXT, INTEGER)
0.093892	(REAL, INTEGER)
0.037091	(TEXT, REAL)
0.033043	(DATE/TIME)
0.019507	(TEXT, REAL, INTEGER)

General Profiling (continued)

Performance Improvement

Small datasets(for all 1825 sets): 3 days to no more than 6 hours

Medium dataset(for one datasets): 3 hours each to 5 minutes each

Large dataset(for one datasets): Half a day each to 30 minutes each

Extra-Large dataset (for one datasets): After improvement, about 4 – 6 hours each

Semantic Profiling

Methodologies

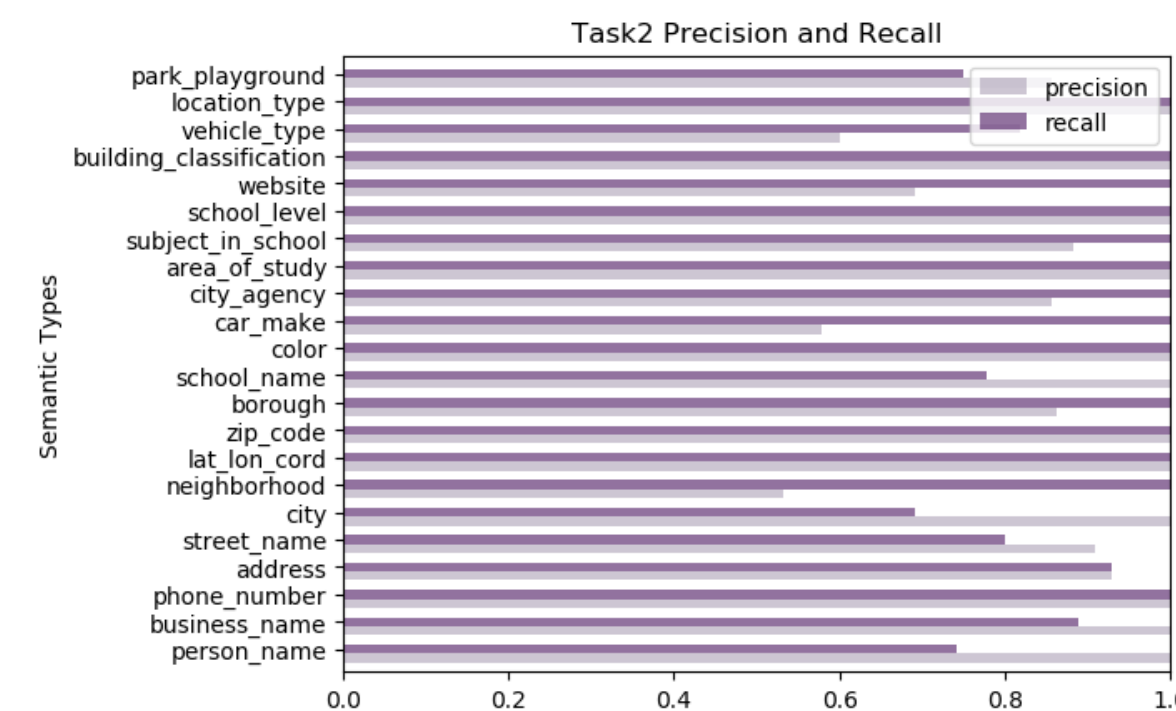
- Data Pre-Processing
 - Manually label all columns
 - One column may have multiple labels
 - Store the labels in a .csv file
 - Extract all distinct values that occur more than a given threshold of times, and group them by labels
 - Extract top 5 frequent values from the metadata json file derived in task1
- Prediction
 - Column names
 - Regular expression
 - Frequent Values (Fast)
 - Works well when there are dominant words in a semantic type
 - Likely to generate false positive
 - Search from known values (Safe, but slow)
 - Best solution if a semantic type has a limited value domain. Such as 'subject_in_school'.
 - External Libraries

Challenges

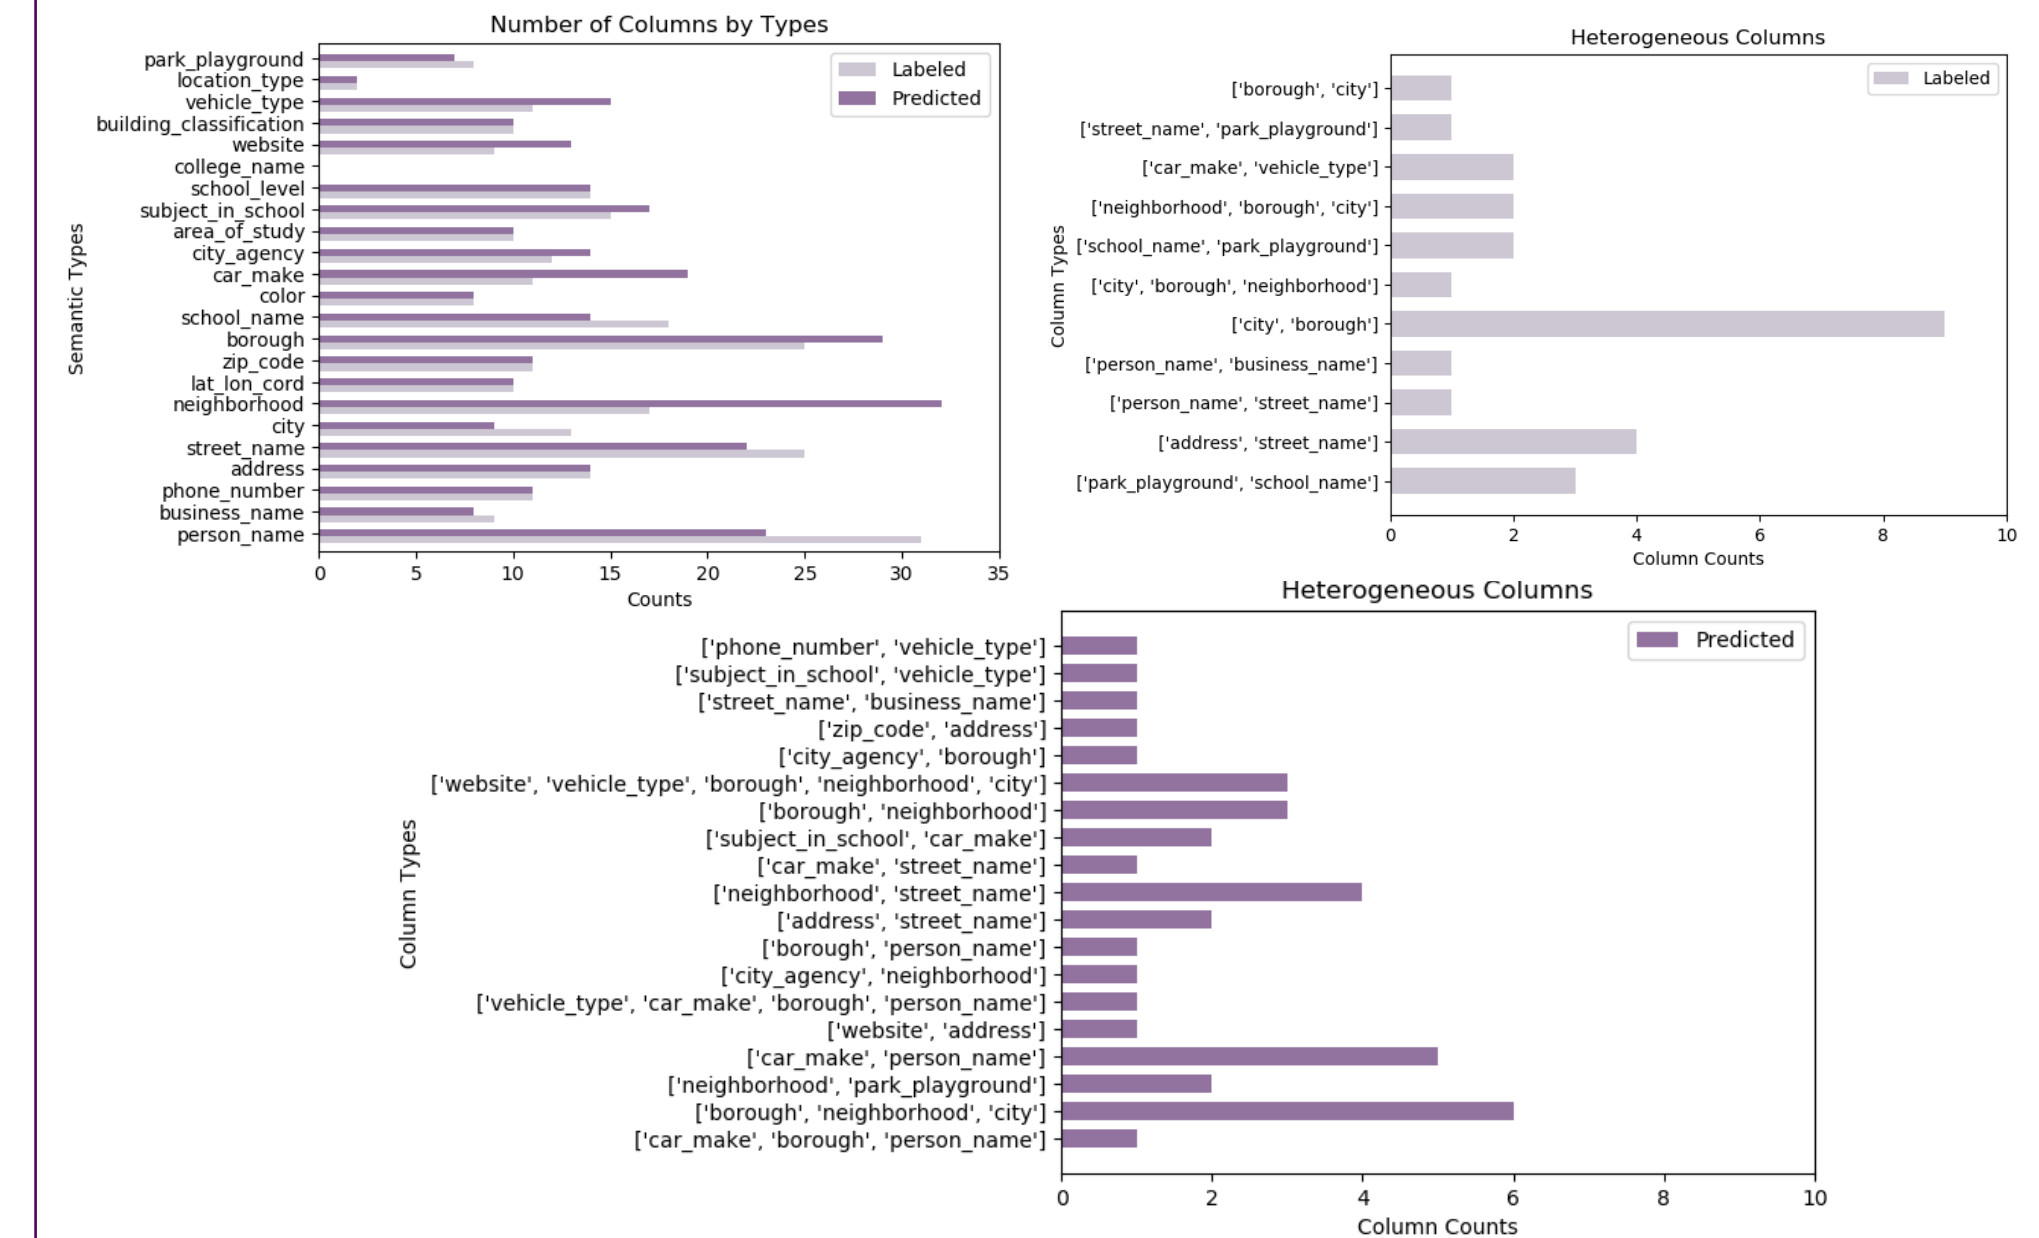
- Unclear Semantic Value
- Short value
- A value can have multiple semantic types
 - Brooklyn can be both 'borough' and 'city'
- Outliers

Semantic Profiling: Results

Semantic Type	Precision	Recall	F1-Score
person_name	1.000000	0.741935	0.851852
business_name	1.000000	0.888889	0.941176
phone_number	1.000000	1.000000	1.000000
address	0.928571	0.928571	0.928571
street_name	0.909091	0.800000	0.851064
city	1.000000	0.692308	0.818182
neighborhood	0.531250	1.000000	0.693878
lat_lon_cord	1.000000	1.000000	1.000000
zip_code	1.000000	1.000000	1.000000
borough	0.862069	1.000000	0.925926
school_name	1.000000	0.777778	0.875000
color	1.000000	1.000000	1.000000
car_make	0.578947	1.000000	0.733333
city_agency	0.857143	1.000000	0.923077
area_of_study	1.000000	1.000000	1.000000
subject_in_school	0.882353	1.000000	0.937500
school_level	1.000000	1.000000	1.000000
website	0.692308	1.000000	0.818182
building_classification	1.000000	1.000000	1.000000
vehicle_type	0.600000	0.818182	0.692308
location_type	1.000000	1.000000	1.000000
park_playground	0.857143	0.750000	0.800000



Semantic Profiling: Results (continued)



Candidate Columns being Keys of a Table

- Naïve Method
 - Select the columns which have N distinct values (N is the number of rows)
 - Can find such columns in more than half datasets
- Brute Force Method
 - Try all possible combinations of columns (only applies to small datasets)
- Improvement(apply to larger datasets)
 - Randomization
 - Sampling
 - Pruning
- Duplicate Rows: Remove or Not

Data Analysis: Null Values and Outliers

Null values

- Finding null values: A combination of observation and auto-detection
- Observation: N/A, unknown, not applicable, repeated short values('!', '-', 's', etc.)
- Use libraries(like Pandas, NumPy) to help detect and analyze

Outliers

- Detect high frequency outliers (utilize task1 results)
- Distance-based outliers

Dealing with Null and Outliers

- Deleting the row / Replace the value with nan
- Use mean to replace the value

Summary

During the process of profiling and analyzing the data, we explored, experimented and gained experience with a variety of tools, technologies and methodologies to deal with Big Data. There is also plenty of work that can be done in the future, like applying our methods to other datasets and try to improve the methods with new results we get.

Contact

Email: wm1065@nyu.edu; yh2834@nyu.edu; xz2456@nyu.edu
GitHub Repository: <https://github.com/wm1065/Big-Data-Final-Project>
Results HDFS directory: /user/yh2834/2019-BigDataResults/

References

- Mining of massive datasets [Rajaraman, A., Leskovec, J., Ullman, J. 2012]
- Profiling relational data: a survey. [Abedjan et al., VLDB 2015]
- Pandas: a Foundational Python Library for Data Analysis and Statistics [W. McKinney, scipy2010]
- Algorithms for Mining Distance-Based Outliers in Large Datasets [Knorrr and Ng, VLDB 1998]